

# A Practical Introduction to Machine Learning in Python

## Day 5 – Friday

### »Next steps«

---

Damian Trilling  
Anne Kroon

d.c.trilling@uva.nl, @damian0604  
a.c.kroon@uva.nl, @annekroon

September 30, 2021

## This part: State of the art and next steps

Hot and happening: Transformers

Do I need all this fancy stuff?

Your takeaway

# Hot and happening: Transformers

---

# The idea

## BERT: Bidirectional Encoder Representations from Transformers (Devlin et al., 2019)

- (Huge) pre-trained model (by, e.g., Google) that is fine-tuned for specific task (by you)
- In simple neural networks, identical words have identical meanings – but meaning can depend on context
- Therefore, the model should take context into accounts. For example, in LSRTM we use *sequences* of words.
- But meaning of a word does not necessarily depend *sequentially* on the preceeding words in that order
- Solution: *Learn* which tokens *to attend to* (attention)

# The idea

## BERT: Bidirectional Encoder Representations from Transformers (Devlin et al., 2019)

- (Huge) pre-trained model (by, e.g., Google) that is fine-tuned for specific task (by you)
- In simple neural networks, identical words have identical meanings – but meaning can depend on context
- Therefore, the model should take context into accounts. For example, in LSRTM we use *sequences* of words.
- But meaning of a word does not necessarily depend *sequentially* on the preceeding words in that order
- Solution: *Learn* which tokens *to attend to* (attention)

# The idea

## BERT: Bidirectional Encoder Representations from Transformers (Devlin et al., 2019)

- (Huge) pre-trained model (by, e.g., Google) that is fine-tuned for specific task (by you)
- In simple neural networks, identical words have identical meanings – but meaning can depend on context
- Therefore, the model should take context into accounts. For example, in LSRTM we use *sequences* of words.
- But meaning of a word does not necessarily depend *sequentially* on the preceeding words in that order
- Solution: *Learn* which tokens *to attend to* (attention)

# The idea

## BERT: Bidirectional Encoder Representations from Transformers (Devlin et al., 2019)

- (Huge) pre-trained model (by, e.g., Google) that is fine-tuned for specific task (by you)
- In simple neural networks, identical words have identical meanings – but meaning can depend on context
- Therefore, the model should take context into accounts. For example, in LSRTM we use *sequences* of words.
- But meaning of a word does not necessarily depend *sequentially* on the preceeding words in that order
- Solution: *Learn which tokens to attend to* (attention)

# The idea

## BERT: Bidirectional Encoder Representations from Transformers (Devlin et al., 2019)

- (Huge) pre-trained model (by, e.g., Google) that is fine-tuned for specific task (by you)
- In simple neural networks, identical words have identical meanings – but meaning can depend on context
- Therefore, the model should take context into accounts. For example, in LSRTM we use *sequences* of words.
- But meaning of a word does not necessarily depend *sequentially* on the preceeding words in that order
- Solution: *Learn* which tokens *to attend to* (attention)



# The idea

## BERT: Bidirectional Encoder Representations from Transformers (Devlin et al., 2019)

- We can use BERT for a lot of different tasks: for sequence-to-sequence predictions (e.g., translation), but also for classification (yeah!)
- Can be done in keras
- The Huggingface transformers library includes a lot of BERT-models (e.g., BERTje for Dutch)

# The idea

## BERT: Bidirectional Encoder Representations from Transformers (Devlin et al., 2019)

- We can use BERT for a lot of different tasks: for sequence-to-sequence predictions (e.g., translation), but also for classification (yeah!)
- Can be done in keras
- The Huggingface transformers library includes a lot of BERT-models (e.g., BERTje for Dutch)

# The idea

## BERT: Bidirectional Encoder Representations from Transformers (Devlin et al., 2019)

- We can use BERT for a lot of different tasks: for sequence-to-sequence predictions (e.g., translation), but also for classification (yeah!)
- Can be done in keras
- The Huggingface transformers library includes a lot of BERT-models (e.g., BERTje for Dutch)

Let's look at an example (imdb.ipynb)

# Hot and happening: Transformers

---

Do I need all this fancy stuff?

# Things to consider

How important is...

- precision/recall? Am I satisfied with .88 when .90 is theoretically possible? .85? .80? .75?
- explainability?
- computational resources?
- generalizability and out-of-sample performance?

# Things to consider

How important is...

- precision/recall? Am I satisfied with .88 when .90 is theoretically possible? .85? .80? .75?
- explainability?
- computational resources?
- generalizability and out-of-sample performance?

# Things to consider

How important is...

- precision/recall? Am I satisfied with .88 when .90 is theoretically possible? .85? .80? .75?
- explainability?
- computational resources?
- generalizability and out-of-sample performance?



# Things to consider

How important is...

- precision/recall? Am I satisfied with .88 when .90 is theoretically possible? .85? .80? .75?
- explainability?
- computational resources?
- generalizability and out-of-sample performance?

# Do I need all this fancy stuff?

- Always estimate a simple baseline model first
- Invest in good hyperparameter-tuning (cross-validation, gridsearch) and don't forget to set aside unseen data for the *final* evaluation.
- If you (a) need to get the highest possible accuracy, or (b) have reasons to assume that the model does not generalize well enough (overfitting problems, bad out-of-sample prediction (e.g., training topics on newspaper 1, predicting topics in newspaper 2)), try embedding-based approaches, transformers, etc.
- Rule of thumb: the more abstract/latent what you want to predict, the less likely classic ML is going to work

## Do I need all this fancy stuff?

- Always estimate a simple baseline model first
- Invest in good hyperparameter-tuning (cross-validation, gridsearch) and don't forget to set aside unseen data for the *final* evaluation.
- If you (a) need to get the highest possible accuracy, or (b) have reasons to assume that the model does not generalize well enough (overfitting problems, bad out-of-sample prediction (e.g., training topics on newspaper 1, predicting topics in newspaper 2)), try embedding-based approaches, transformers, etc.
- Rule of thumb: the more abstract/latent what you want to predict, the less likely classic ML is going to work

## Do I need all this fancy stuff?

- Always estimate a simple baseline model first
- Invest in good hyperparameter-tuning (cross-validation, gridsearch) and don't forget to set aside unseen data for the *final* evaluation.
- If you (a) need to get the highest possible accuracy, or (b) have reasons to assume that the model does not generalize well enough (overfitting problems, bad out-of-sample prediction (e.g., training topics on newspaper 1, predicting topics in newspaper 2)), try embedding-based approaches, transformers, etc.
- Rule of thumb: the more abstract/latent what you want to predict, the less likely classic ML is going to work


## Do I need all this fancy stuff?

- Always estimate a simple baseline model first
- Invest in good hyperparameter-tuning (cross-validation, gridsearch) and don't forget to set aside unseen data for the *final* evaluation.
- If you (a) need to get the highest possible accuracy, or (b) have reasons to assume that the model does not generalize well enough (overfitting problems, bad out-of-sample prediction (e.g., training topics on newspaper 1, predicting topics in newspaper 2)), try embedding-based approaches, transformers, etc.
- Rule of thumb: the more abstract/latent what you want to predict, the less likely classic ML is going to work

## Your takeaway

---

(short recap of course)



*Have your plans about how to and  
wether to use ML changed?*





*What are your next steps?*

Last part: we help you working on (or discussing about) your own projects.

# References

---



Devlin, J., Chang, M. W., Lee, K., & Toutanova, K. (2019). BERT: Pre-training of deep bidirectional transformers for language understanding. In *Naacl hlt 2019 - 2019 conference of the north american chapter of the association for computational linguistics: Human language technologies - proceedings of the conference*.